
Model Inference from Neural Artifacts: An Exchange Law for Structure and Memory

Anonymous Author(s)

Abstract

1 An artifact’s inference state has two capacities, cardinality and chain height, and
2 independent pooling is priced by the second. The state is the set of distinctions that
3 future observations can expose, a finite meet-semilattice $\mathcal{R}$; it is what an artifact
4 is obliged to retain, where weight-space learning estimates statistics of what it
5 contains. Coupled attention needs $\lceil \log_b |\mathcal{R}| \rceil$ coordinates of b values; independently
6 pooled channels of sizes $a_1, \dots, a_d$ need $\sum_j (a_j - 1) \geq \text{ht}(\mathcal{R}) - 1$, for arbitrary
7 commutative pooling and redundant implementation states. On an n -state chain
8 the two are exact and separate exponentially, $\lceil (n - 1)/(b - 1) \rceil$ against $\lceil \log_b n \rceil$;
9 on a Boolean store of k independently eliminated candidates they coincide at k .
10 Chain height is the number of times future evidence can change the answer, so on
11 a pooled memory adaptive queries and memory are one resource: a policy whose
12 memory is the store makes at most $\sum_j (a_j - 1)$ informative queries, whereas a
13 decoder with unrestricted memory exchanges stored bits and query bits one for one.
14 Attention–ReLU artifacts realize both sides of each law, their minimal inference
15 machines are recovered from operational parameters, and finite-symbol hardening
16 preserves the machine whenever the probability mass on wrong symbols lies below
17 an explicit threshold. All results in this paper are verified in Lean 4.¹

18 1 Introduction

19 Neural artifacts contain more than the answers returned on an evaluation set. Their weights, routing
20 decisions, and intermediate representations also determine how information is stored, combined, and
21 used after new evidence arrives. Learning directly on network parameters already supports model
22 comparison and prediction, with parameter symmetries providing a central architectural principle
23 [22]. We study a complementary question: *which model of inference can be recovered from an*
24 *artifact, and which distinctions must its internal representations preserve?*

25 A network can represent candidate models through different mechanisms. Attention can expose a
26 selection rule or a compositional routing tree; feedforward computations can retain tables, transforma-
27 tions, and residual information. These roles do not assign an intrinsic meaning to a layer type. Instead,
28 we specify their common semantic interface: each representation participates in maintaining the
29 model possibilities consistent with observations. The central issue is then whether information stored
30 in one representation can replace information stored in another while preserving future inference.

31 Counting states gives a necessary condition, but misses the cost of composition. A four-state chain
32 can be encoded by two bits. It cannot be maintained by two independently pooled binary channels,
33 even when their pooling operations are chosen freely. Three such channels suffice. The discrepancy is
34 not a shortage of bits at a single instant: it arises because each channel must combine evidence without
35 consulting the others. We prove the exact law for every chain size and every finite commutative
36 pooling architecture. At fixed state cardinality, a Boolean store has a different exchange law because
37 complementary coordinates preserve intersection independently.

¹Lean sources, paper source, and checks: <https://anonymous.4open.science/r/neural-artifacts-exchange-law-838F/>

Our semantic starting point is a version space [12]. An observation removes incompatible candidates, and an inference rule reads the survivors. We identify two version spaces when no permitted continuation changes the inference rule’s answer between them. This is the classical automata-minimization construction [13, 16], applied to the meet-semilattice generated by model observations. Every correct implementation has a unique interpretation onto this quotient on its reachable states. The inference task forces the semantic object.

The inference state $\mathcal{R}$ forced by a task carries two numbers, $|\mathcal{R}|$ and $\text{ht}(\mathcal{R})$. Coupled pooling is priced by the first and independent pooling by the second, so the exchange rate between structure and memory is a property of the store, and the separation of corollary 4 appears exactly when $\text{ht}(\mathcal{R})$ approaches $|\mathcal{R}|$. Chain height is mind-change complexity, the number of times a readout can change along a history; a store that affords many revisions at logarithmic width must couple its coordinates. The same quantity bounds adaptive experimentation on a pooled memory (theorem 11), which places the memory law of appendix G and the exchange law on one scale. Finite-state extraction and an explicit attention–ReLU compiler connect the laws to artifacts, and finite-symbol hardening preserves the extracted machine, including ties between branches with the same symbol.

Position relative to prior work. Knowledge compilation compares representations through their succinctness and supported queries and transformations [5]. We apply this perspective to the persistent state of neural inference and derive sharp costs for independently compositional representations. The kernel criterion for exact factorization is the classical subdirect-product criterion [2]; the new quantitative result is the exact exchange region for pooled inference, together with matching neural constructions. Automata extraction from recurrent networks uses queries and abstractions to recover state dynamics [20]. Our extraction theorem supplies an exact bridge from a finite neural execution interface to the same canonical object used by the lower bounds. Soft attention can simulate hard attention under score-gap conditions and preserves a final classifier [21]; finite-symbol correction preserves the entire model-inference machine. Slot Attention provides an important example of exchanging information between perceptual features and object-level representations [11]. The algebraic interface below specifies what must be verified to give such an exchange an exact model-inference semantics. Chain height is the mind-change complexity of inductive inference [3, 6]; the exchange law prices it by the pooling operation. Weight-space learning treats the artifact as data: equivariant architectures on parameter space [15], graph metanetworks over computation graphs [9, 10], accuracy predicted from weights [19], populations of trained models [18], and an atlas of public models [7]. These learn statistics of the artifact; we characterize what the artifact is obliged to retain.

2 The inference state represented by an artifact

Let $\mathcal{C}$ be a finite set of candidate models, P an observation alphabet, and $K : P \rightarrow \mathcal{P}(\mathcal{C})$ a consistency map. The alphabet can include labeled examples, experimental results, or inspected artifact properties. For a history $s = p_1 \cdots p_t$, define

$$V(s) = \bigcap_{i=1}^t K(p_i), \quad V(\epsilon) = \mathcal{C}, \quad \mathcal{L} = \{V(s) : s \in P^*\}. \quad (1)$$

The finite family $\mathcal{L}$ contains $\mathcal{C}$ and is closed under intersection, since $V(st) = V(s) \cap V(t)$. Fix a readout $\rho : \mathcal{L} \rightarrow Y$. First-consistent-model selection and query certification are examples. Empty version spaces receive an explicit inconsistency answer whenever consistency is relevant.

Definition 1 (Continuation equivalence). For $U, V \in \mathcal{L}$, let

$$U \equiv_\rho V \iff \rho(U \cap H) = \rho(V \cap H) \text{ for every } H \in \mathcal{L}. \quad (2)$$

Write $\mathcal{R} = \mathcal{L}/\equiv_\rho$, with $[U] \wedge [V] = [U \cap V]$ and top element $[\mathcal{C}]$.

The relation respects intersection: $H \cap W \in \mathcal{L}$, so equivalence of U, V implies equivalence of $U \cap W, V \cap W$. Taking $H = \mathcal{C}$ shows that the readout descends to $\mathcal{R}$.

Proposition 2 (Canonical inference state, after Moore [13] and Nerode [16]). *The least number of states of a deterministic machine computing $s \mapsto \rho(V(s))$ is $|\mathcal{R}|$. The same minimum holds for machines that pool observations in a finite commutative monoid. Every correct deterministic implementation has a unique map δ from its reachable states onto $\mathcal{R}$ satisfying*

$$\delta(m_\epsilon) = [\mathcal{C}], \quad \delta(\text{step}(m, p)) = \delta(m) \wedge [K(p)], \quad \text{out}(m) = \bar{\rho}(\delta(m)). \quad (3)$$

86 *Proof.* For a state reached by s , set $\delta(m_s) = [V(s)]$. If $m_s = m_t$, every common continuation
 87 produces the same output. Every $H \in \mathcal{L}$ is the version space of a continuation, so $V(s) \equiv_\rho V(t)$.
 88 This proves well-definedness. Every quotient class is reachable; the update and readout equations
 89 follow from intersection. Induction on histories proves uniqueness. Hence any implementation has
 90 at least $|\mathcal{R}|$ reachable states. Conversely, initialize at $[\mathcal{C}]$, pool $[K(p)]$ by meet, and apply $\bar{\rho}$. This
 91 commutative machine has exactly $|\mathcal{R}|$ states. $\square$

92 The quotient depends on both the available observations and the required answers. A constant readout
 93 gives one state even when $\mathcal{L}$ is large. An injective readout gives $\mathcal{R} = \mathcal{L}$. Enlarging the observation
 94 language or enriching the readout changes the distinctions that an artifact must retain. In particular, a
 95 width- W state with b possible values per coordinate satisfies

$$|\mathcal{R}| \leq b^W. \quad (4)$$

96 Here b is an alphabet size: p -bit coordinates have $b = 2^p$. Equation (4) counts persistent inference
 97 states, independently of the implementation of the candidate models.

98 The second number is the height $\text{ht}(\mathcal{R})$, the number of elements in a longest chain. Along any history
 99 the states $[V(s)]$ descend in $\mathcal{R}$, so the readout changes at most $\text{ht}(\mathcal{R}) - 1$ times. Cardinality is what
 100 a coupled store must hold; height is what an independent one must hold (section 3).

101 **Concrete stores.** Nested consistency sets $U_0 \subsetneq \dots \subsetneq U_{n-1} = \mathcal{C}$, together with a readout that
 102 distinguishes them, give the chain C_n . Positive observations of ordered thresholds give this form
 103 with a suitable domain and readout. Independent candidate eliminations generate the Boolean store
 104 $2^{[k]}$; reading the full survivor set distinguishes all 2^k states.

105 3 An exact exchange law

106 An independent pooled channel consists of a finite commutative monoid M_j and an observation
 107 embedding $e_j : P \rightarrow M_j$. Its state after s is

$$\mu_j(s) = \prod_{p \in s} e_j(p). \quad (5)$$

108 A joint readout can inspect every channel, but each channel combines observations using its own
 109 operation. The monoids need not be idempotent, and different histories with the same inference state
 110 can have different implementations. We write $\text{ht}(Q)$ for the number of elements in a longest chain of
 111 a finite partially ordered set Q .

112 **Theorem 3** (Exchange under independent pooling). *Suppose d independent pooled channels compute*
 113 *the inference task of section 2. Let $E_j = \text{Idem}(M_j) = \{x : x^2 = x\}$, ordered by $x \leq y$ when*
 114 *$xy = x$. Then*

$$|\mathcal{R}| \leq \prod_{j=1}^d |M_j|, \quad \text{ht}(\mathcal{R}) - 1 \leq \sum_{j=1}^d (\text{ht}(E_j) - 1) \leq \sum_{j=1}^d (|M_j| - 1). \quad (6)$$

115 For $\mathcal{R} = C_n$ and channel budgets $a_j \geq 1$, a correct implementation with $|M_j| \leq a_j$ exists if and
 116 only if

$$\sum_{j=1}^d (a_j - 1) \geq n - 1. \quad (7)$$

117 *Proof.* Let S be the image of $s \mapsto (\mu_1(s), \dots, \mu_d(s))$. It is a submonoid of $\prod_j M_j$. By proposi-
 118 tion 2, the map $\delta : S \rightarrow \mathcal{R}$ is well-defined and surjective. It is a monoid homomorphism because
 119 concatenation corresponds to multiplication in S and meet in $\mathcal{R}$.

120 Every element of a finite monoid has an idempotent positive power. Since $\delta(x^r) = \delta(x)$ for $r \geq 1$, the
 121 restriction $\delta : \text{Idem}(S) \rightarrow \mathcal{R}$ remains surjective. Commutativity makes $\text{Idem}(S)$ a meet-semilattice,
 122 contained in $\prod_j E_j$. Lift a chain $r_1 < \dots < r_h$ in $\mathcal{R}$ to idempotents e_i and set $z_i = e_i \cdots e_h$.
 123 Then $z_i \leq z_{i+1}$ and $\delta(z_i) = r_i$, so the lifted chain is strict. A chain in a product has at most

124 $1 + \sum_j (\text{ht}(E_j) - 1)$ elements: every strict step advances at least one coordinate. This proves (6)
 125 and the necessity of (7).

126 For sufficiency, label the chain $0 < \dots < n - 1$. Distribute its $n - 1$ adjacent boundaries among sets
 127 B_j , with $|B_j| \leq a_j - 1$. The map

$$f_j(r) = |\{c \in B_j : c < r\}| \quad (8)$$

128 preserves minimum and takes $|B_j| + 1$ values. Every boundary is distinguished by some coordinate,
 129 so the paired map is injective. Pool with minimum in each coordinate and decode the joint state. $\square$

130 The use of idempotent powers is essential: the lower bound applies to redundant encodings, cyclic
 131 state components, and all other finite commutative pooling operations. In particular, a finite group
 132 contributes no order height because its identity is its only idempotent.

133 **Corollary 4** (Independent versus coupled width). *For an n -state chain and $b \geq 2$, the minimum*
 134 *number of independently pooled b -value coordinates is*

$$W_{\text{ind}}(n, b) = \left\lceil \frac{n - 1}{b - 1} \right\rceil. \quad (9)$$

135 *If the pooling operation can mix coordinates, the exact minimum is*

$$W_{\text{joint}}(n, b) = \lceil \log_b n \rceil. \quad (10)$$

136 *Proof.* The first equality is theorem 3. For the second, encode the rank in base b and implement
 137 minimum on the represented ranks. The cardinality bound (4) is necessary and attained. Place unused
 138 codes below the reachable chain to extend minimum to the full b^W -element carrier. The encoded top
 139 remains the monoid identity and unused codes are not reached. $\square$

140 For $n = 2^k$, independent binary pooling requires $2^k - 1$ persistent coordinates, while coupled pooling
 141 requires k . A coupled head can implement minimum directly: assign keys $-r, -s$, query 1, and
 142 values equal to the digit encodings of r, s . Hard attention selects the smaller rank. For equal ranks
 143 the values coincide, so any tie rule gives the same result.

144 **Corollary 5** (Width for an arbitrary inference state). *For any finite $\mathcal{R}$ and $b \geq 2$,*

$$W_{\text{joint}}(\mathcal{R}, b) = \lceil \log_b |\mathcal{R}| \rceil, \quad W_{\text{ind}}(\mathcal{R}, b) \geq \max \left\{ \lceil \log_b |\mathcal{R}| \rceil, \left\lceil \frac{\text{ht}(\mathcal{R}) - 1}{b - 1} \right\rceil \right\}. \quad (11)$$

145 *For the Boolean store $2^{[k]}$ with survivor readout, $W_{\text{ind}} = W_{\text{joint}} = k$.*

146 *Proof.* The coupled construction of corollary 4 extends to any $\mathcal{R}$ by sending unused codes to an
 147 adjoined absorbing bottom, which keeps the carrier a commutative idempotent monoid with identity
 148 $[\mathcal{C}]$; (4) forbids fewer coordinates. The lower bound is (4) together with (6) at $|M_j| \leq b$. One binary
 149 channel per candidate, pooled by conjunction, reaches every survivor set of $2^{[k]}$ with an injective
 150 readout, and $|2^{[k]}| = 2^k$ forbids fewer coordinates under any pooling. $\square$

151 The separation belongs to chains: $W_{\text{ind}}/W_{\text{joint}}$ is controlled by $\text{ht}(\mathcal{R})/\log_b |\mathcal{R}|$.

152 3.1 How a structural view replaces residual memory

153 A normalized compositional view of $\mathcal{R}$ is a meet homomorphism $a : \mathcal{R} \rightarrow A$. Its kernel α is a meet
 154 congruence: equivalent states remain equivalent after intersection with any state. A second view
 155 $g : \mathcal{R} \rightarrow G$, with kernel β , jointly determines the inference state precisely when

$$\alpha \cap \beta = \Delta_{\mathcal{R}}. \quad (12)$$

156 This is the subdirect-product criterion [2]. At a fixed structural view, its exact compositional residual
 157 requirement is therefore

$$M_{\text{comp}}(\alpha) = \min_{\substack{\beta \in \text{Con}_{\wedge}(\mathcal{R}) \\ \alpha \cap \beta = \Delta_{\mathcal{R}}}} |\mathcal{R}/\beta|. \quad (13)$$




Figure 1: Exact structure–memory exchange for two 16-state inference tasks. A four-state structural view leaves 13 compositional residual states in a chain and four in a Boolean store with two exposed coordinates. The Boolean line joins the admissible coordinate projections. State cardinality alone does not determine the exchange rate.

Refining the structural view from α to $\alpha' \subseteq \alpha$ enlarges the feasible set in (13); hence $M_{\text{comp}}(\alpha') \leq M_{\text{comp}}(\alpha)$. In contrast, an arbitrary residual code used to reconstruct the state once requires exactly

$$M_{\text{stat}}(\alpha) = \max_{F \in \mathcal{R}/\alpha} |F|. \quad (14)$$

Assign distinct labels within each fibre and reuse them across fibres.

Proposition 6 (Sharp chain and Boolean exchange). *For a chain C_n and any structural meet quotient with a states,*

$$M_{\text{comp}}(\alpha) = n - a + 1. \quad (15)$$

The same minimum holds when the residual is an arbitrary finite commutative pooled channel, including redundant histories. For $\mathcal{R} = 2^{[k]}$ with structural view $U \mapsto U \cap J$, $|J| = s$,

$$M_{\text{comp}}(\alpha) = M_{\text{stat}}(\alpha) = 2^{k-s}. \quad (16)$$

Proof. A chain congruence has interval blocks: if $x \sim z$ and $x \leq y \leq z$, meeting with y gives $x \sim y$. Its a blocks distinguish $a - 1$ boundaries. The residual must distinguish the other $n - a$ boundaries, and cutting exactly these boundaries attains $n - a + 1$ states. The general lower bound follows from theorem 3, with the fixed structural quotient as the first channel. For the Boolean store, every structural fibre has 2^{k-s} elements, requiring that many distinct residual values. Projection to $[k] \setminus J$ attains the bound and preserves intersection. $\square$

For the four-state chain, take structural blocks $\{0, 1\}, \{2, 3\}$. A static binary residual can record parity. It cannot update by itself: 0 and 2 have the same parity, but meeting each with 1 gives 0 and 1. A compositional residual instead uses $\{0\}, \{1, 2\}, \{3\}$. This example separates recoverable information from independently reusable information.

4 Recovering models from neural artifacts

The exchange law concerns the semantic state represented by an artifact. We now give operational extraction and realization procedures. An artifact specifies its architecture, parameters, initial state, numerical semantics, and readout; these determine a transition function $F_\theta(z, p)$.

Theorem 7 (Finite-state extraction). *Let Z be a finite state codebook containing z_0 , let P be finite, and suppose $F_\theta : Z \times P \rightarrow Z$ and $o_\theta : Z \rightarrow Y$ can be evaluated exactly. From the operational artifact one can recover the reachable transition table, its minimal Moore machine, and a state-by-state simulation onto that machine. If the artifact computes $\rho(V(s))$ for every $s \in P^*$, the extracted minimal machine is isomorphic to $\mathcal{R}$. If its minimized input transitions commute and are idempotent, the artifact itself determines its meet-semilattice representation, and any meet law consistent with its transitions is that one.*

Proof. Starting at z_0 , enumerate every transition from each newly reached state. This terminates after at most $|Z||P|$ evaluations. Partition the reachable states by readout, and repeatedly refine by successor-block signatures until stable. Induction on continuation length shows that the final equivalence is agreement on every future output. The quotient and projection give the required

190 machine and simulation. Proposition 2 identifies this quotient with $\mathcal{R}$. For the final statement, choose
 191 an access word w_r for each minimal state r , and define $r \wedge s$ by running w_s from r . Commutativity
 192 makes this independent of the access word; idempotence makes it a meet. The initial state is its top.
 193 Appendix B proves the representation and its uniqueness. $\square$

194 A finite-precision execution supplies a finite codebook directly, with every persistent register included.
 195 The extraction is invariant under function-preserving hidden-neuron permutations because it depends
 196 on the transition function, not the ordering of its implementation units. State-coordinate relabelings
 197 yield isomorphic extracted machines. These are the semantic symmetries relevant to model inference;
 198 parameter-level equivariance [15, 22] respects the same permutations without identifying what they
 199 preserve.

200 The last clause gives a fully operational analysis: extract the transition system, minimize it, test
 201 commutativity and idempotence of its input transitions, and recover the meet operation. The candidate
 202 realization $\mathcal{C} = \mathcal{R}$, $K(p) = \{r : r \leq t_p(\top)\}$ then has survivor sets $\downarrow r$. The artifact determines
 203 this operational version-space model; an external model population supplies its application-specific
 204 interpretation. For continuous-state artifacts, a finite invariant-region certificate gives the same extrac-
 205 tion theorem; rational hard-attention/ReLU maps and semialgebraic regions admit exact certificate
 206 checking (appendix B).

207 4.1 Attention forests and finite feedforward maps

208 We use a typed code grammar with finite input alphabet X . A leaf is a finite map $f : X \rightarrow Y$. An
 209 internal node has children $c_1, \dots, c_k$ with common finite output alphabet Y , affine scores $s_j(x)$ in a
 210 declared feature vector $\phi(x)$, and a finite map $g : Y \rightarrow Y'$. Its interpretation is

$$\llbracket c \rrbracket(x) = g(\llbracket c_{j^*(x)} \rrbracket(x)), \quad j^*(x) = \min \arg \max_j s_j(x). \quad (17)$$

211 A forest provides multiple such outputs. Fixed attention reads followed by finite tuple maps compose
 212 these outputs into a finite-state update or a model-query program. Neither the grammar nor the store
 213 requires the candidate models themselves to be linear.

214 **Theorem 8** (Neural realization and extraction). *Every finite attention forest of the form (17) has*
 215 *a masked dot-product attention implementation with shared positionwise ReLU feedforward maps,*
 216 *residual connections, and no normalization. Its operational parameters determine an extracted code*
 217 *satisfying*

$$\llbracket \text{extract}(\text{compile}(c)) \rrbracket(x) = \llbracket c \rrbracket(x) \quad \text{for every } x \in X. \quad (18)$$

218 *Any finite transition system can be realized by a ReLU update, and composing either construction*
 219 *with theorem 7 recovers its minimal inference machine.*

220 *Construction.* Encode a symbol y as the basis vector e_y . A finite map has matrix $H_g e_y = e_{g(y)}$. An
 221 affine score $\alpha_j^\top \phi(x) + \beta_j$ is a dot product between query $(\phi(x), 1)$ and key (α_j, β_j) . Hard attention
 222 selects the child value, and the feedforward map applies H_g . For shared feedforward dispatch, pair
 223 a node-identity coordinate u_i and a symbol coordinate v_j using $\text{ReLU}(u_i + v_j - 1)$; this is their
 224 conjunction on one-hot inputs. Masks isolate each parent’s children. Separate protected, payload,
 225 and scratch coordinates preserve features and implement residual overwrites. Appendix C gives the
 226 complete layer construction and the inductive proof.

227 For a transition table $T : [N] \times [m] \rightarrow [N]$, the hidden units

$$h_{ij}(z, p) = \text{ReLU}(z_i + p_j - 1), \quad F(z, p) = \sum_{i,j} e_{T(i,j)} h_{ij}(z, p) \quad (19)$$

228 implement $F(e_i, e_j) = e_{T(i,j)}$ exactly. Extraction evaluates these operational maps and reads the
 229 routing scores and wiring. No unused source-code identifier is required. $\square$

230 The compiler realizes both sides of corollary 4. Independent channels implement the cut maps (8); a
 231 joint attention operation implements minimum on compact digit codes. Compositional structure sup-
 232 ports model inference when the extracted maps preserve the relevant operations; a large feedforward
 233 table supports the same task by retaining the corresponding transition information.

5 Hardening preserves inference through finite symbols

Changing softmax weights to argmax decisions changes the executed function in general. We instead give an exact preservation condition at the symbolic interfaces of theorem 8. The correction function below is the standard ReLU rounding construction used in soft-attention simulations [21], parameterized by a tolerance $0 < \eta < 1/2$:

$$R_\eta(t) = \frac{\text{ReLU}(t - \eta) - \text{ReLU}(t - 1 + \eta)}{1 - 2\eta}. \quad (20)$$

Apply it coordinatewise before the finite map H_g .

Theorem 9 (Exact symbolic hardening). *At a gate with scores s_j and one-hot child values e_{y_j} , let j^* be the deterministic hard winner and $y^* = y_{j^*}$. For temperature $T > 0$, write*

$$p_j(T) = \frac{e^{s_j/T}}{\sum_i e^{s_i/T}}, \quad \varepsilon_{\neq y^*}(T) = \sum_{j: y_j \neq y^*} p_j(T). \quad (21)$$

Then

$$R_\eta \left(\sum_j p_j(T) e_{y_j} \right) = e_{y^*} \iff \varepsilon_{\neq y^*}(T) \leq \eta. \quad (22)$$

Suppose every maximum-score branch returns y^* , there are m_* such branches, and every branch with a different symbol has score at most $s_{j^*} - \delta$ for some $\delta > 0$. If k_- branches have a different symbol, then

$$\varepsilon_{\neq y^*}(T) \leq \frac{r(T)}{1 + r(T)}, \quad r(T) = \frac{k_-}{m_*} e^{-\delta/T}. \quad (23)$$

A forest whose every gate satisfies (22) on its symbolic inputs computes exactly the same function at temperature T and at hard attention. When used as a state update, it preserves every history and hence the extracted inference quotient and all exchange laws above.

Proof. The coordinate of y^* in the mixture is $1 - \varepsilon_{\neq y^*}$, and the other coordinates sum to $\varepsilon_{\neq y^*}$. The map R_η is zero on $[0, \eta]$, one on $[1 - \eta, 1]$, and strictly between them otherwise. This proves (22). The total unnormalized weight of different symbols is at most $k_- e^{(s_{j^*} - \delta)/T}$, while the weight of y^* is at least $m_* e^{s_{j^*}/T}$. Their ratio proves (23). Induction over the forest restores an exact symbol after each gate, so errors do not accumulate with depth. Induction over state updates then proves equality on all histories. $\square$

For a unique winner with gap δ from every other branch, a sufficient condition is

$$T \leq \frac{\delta}{\log((k - 1)(1 - \eta)/\eta)} \quad (k \geq 2). \quad (24)$$

The symbol-based criterion is stronger than unique-route stability: tied or changing routes cause no loss when they return the same symbol. Conversely, branches with different values at an unresolved tie can invalidate hardening. Without correction, a soft gate mixing constants 0, 1 returns an interior value that a hard selection of those experts does not return.

Execution and identification are different. Let $O_T(\theta)$ expose routing probabilities and let $O_0(\theta) = h(O_T(\theta))$ expose only their winners, with other observed quantities held fixed. The observer's version spaces satisfy

$$\{\theta' : O_T(\theta') = O_T(\theta)\} \subseteq \{\theta' : O_0(\theta') = O_0(\theta)\}. \quad (25)$$

Thus hardening can preserve execution while a coarsened trace reveals less about the artifact. More generally, a target property $Q(\theta)$ can be inferred exactly from behavioral observations $O(\theta)$ and an artifact statistic $A(\theta)$ if and only if

$$\ker(O, A) \subseteq \ker Q. \quad (26)$$

Artifact access improves identification precisely when it separates a target-relevant observational equivalence class. Which equivalence class is separated depends on the observation channel.

6 Model queries, memory, and interventions

A query $q : \mathcal{C} \rightarrow \mathcal{A}$ can be certified before a unique model is identified. For nonempty V , return a when every $c \in V$ has $q(c) = a$, and undetermined otherwise. Return inconsistent on $V = \emptyset$. This defines another readout in section 2 and therefore another canonical quotient.

Proposition 10 (Exact query certificate). *Fix a true model $c^* \in V$ and deterministic experiments with outcomes $D(c, e)$. Let*

$$B_q = \{c \in V : q(c) \neq q(c^*)\}, \quad A_e = \{c \in B_q : D(c, e) \neq D(c^*, e)\}. \quad (27)$$

Truthful outcomes for an experiment set E certify $q(c^)$ exactly when*

$$\bigcup_{e \in E} A_e = B_q. \quad (28)$$

Consequently the minimum target-relative certificate size is the minimum set-cover size; fixed nonnegative experiment costs give weighted set cover. Stored outcomes replace B_q by its uncovered subset before further experiments are chosen.

Proof. The query-disagreeing survivors are exactly $B_q \setminus \bigcup_{e \in E} A_e$. The true model survives, so the query is certified if and only if this set is empty. $\square$

An artifact can store an experimental outcome, a structured model that entails it, or the update rule that obtains its consequence; these representations exchange information according to the distinctions they remove and the operations they preserve. Sound transfer between two representations is the reduced-product construction of abstract interpretation [4]; appendix E gives its version-space realization.

Theorem 11 (Adaptive experiments on a pooled memory). *Let a policy choose each experiment and its final answer as functions of a persistent state, and let that state be the pooled outcome $\mu(s)$ of (5). Then the inference state $\delta(\mu(s))$ depends only on the set of outcomes received, and the number of experiments that change it along any branch is at most*

$$\text{ht}(\mathcal{R}) - 1 \leq \sum_{j=1}^d (\text{ht}(E_j) - 1). \quad (29)$$

On a coupled W -coordinate b -value memory the count is at most $b^W - 1$. Both bounds are attained.

Proof. By proposition 2, $\delta(\mu(s)) = [V(s)]$, and $V(s)$ is an intersection over the outcome set. Successive states form a descending chain in $\mathcal{R}$, whose strict steps number at most $\text{ht}(\mathcal{R}) - 1$; theorem 3 bounds this by the channel heights, and (4) bounds it by $b^W - 1$. The chain constructions of corollary 4, read as policies that eliminate one boundary per experiment, attain both. $\square$

Finite acyclic structural causal models give an explicit intervention instance [17]. Write $x_i = f_i(x_{\text{pa}(i)}, u_i)$ in topological order. Each f_i is a finite feedforward map, and $\text{do}(x_J = a_J)$ replaces the selected equations by constants. Compiling these maps and replacements gives an attention-ReLU program whose intervention answers agree with the structural equations by induction over the graph. Proposition 10 then characterizes exactly which stored or newly observed intervention outcomes determine a query in a finite model population. Appendix F supplies the construction and a separation between observational agreement and intervention disagreement.

A decoder with unrestricted memory obeys instead $B + t \log_2 q \geq N \log_2 q$ for tables $f : [N] \rightarrow [q]$ reconstructed from a B -bit target-dependent artifact and at most t adaptive queries (appendix G): stored and queried bits exchange one for one because the transcript is retained. A pooled memory retains only its state, so an experiment counts when it moves the state, and memory rungs and informative rounds are the same quantity. The bit law is the coupled case with unrestricted memory. The observation language and model grammar determine which distinctions require table memory: the AES S-box is specified by finite-field inversion and an affine transformation [14], and that grammar, not the table, is the population.

Table 1: Exhaustive execution of the constructions. Counts are enumerated objects or input cases.

Construction	Cases	Verified property
Meet-closed stores with readouts	2,120	Direct and iterative quotients agree
Neural transition-table extraction	72,911	Every transition recovered exactly
Commutative monoids, sizes 1–4	106	Idempotent lifts; 329 chain quotients
Chain structural quotients, sizes 1–7	127	Exact residual minimum
Compiled attention forests	356	All 256 Boolean maps on eight inputs realized
Hard code-extraction input cases	2,848	Exact semantic round trip
Soft code-execution input cases	8,544	Maximum absolute error 0
Shared-block transformer executions	11,392	Exact outputs and state invariants
Symbol-level hardening cases	1,080	Exact criterion, including 210 same-symbol tie cases

7 What an artifact is obliged to retain

Learning on artifacts has a representation-independent target: the inference quotient, its decompositions, and the maps between them. The quotient is what two artifacts share when they implement the same inference; the exchange law separates architectures of equal state cardinality; the hardening certificate names the transformations that preserve all future inference. In a model population [18] the extracted object answers which observations separate models, which queries are already determined, and which experiments change the answer. All results in this paper are verified in Lean 4, and the constructions are executed exhaustively on the finite families of table 1, each enumeration defined in appendix I.

Each transition artifact undergoes a hidden-unit permutation before extraction, and all 11,392 shared-block runs preserve protected coordinates and exact symbolic payloads; appendix I defines each enumeration and its assertions.

Certified artifact transformations. The extracted machine supports three certified operations: state compression, which merges states with the same continuation behavior; representation factorization, which chooses a structural congruence α and computes its minimum compatible residual through (13); and hardening, which preserves the entire transition system through the local symbol-mass certificate. Each is invariant under relabeling of the extracted states, so artifacts are compared by the inference operations they implement and the distinctions those operations retain.

Limitations. The bounds measure persistent inference states, not parameter count or running time. Finite-state extraction has exponential worst-case cost in encoded width, and the computational study verifies compiled constructions; their prevalence under optimization is a separate question. Intervention certificates are relative to the declared model population and experiment semantics. The inference theorems require correctness over the specified observation language; obtaining such certificates for learned continuous-state representations is the next empirical and verification problem.

Conclusion. An artifact’s inference capacity is its retained distinctions together with its composition law: cardinality for a coupled store, chain height for an independent one, and the long revision budget at logarithmic width is the capacity in-context evidence draws on. Attention and feedforward computation realize both laws, finite-symbol hardening preserves their semantic state, and model extraction, representation exchange, and artifact comparison become one operational problem.

References

- [1] Saugata Basu, Richard Pollack, and Marie-Françoise Roy. *Algorithms in Real Algebraic Geometry*. Springer, second edition, 2006. doi: 10.1007/3-540-33099-2.
- [2] Garrett Birkhoff. Subdirect unions in universal algebra. *Bulletin of the American Mathematical Society*, 50:764–768, 1944. URL <https://projecteuclid.org/journals/bulletin-of-the-american-mathematical-society/volume-50/issue-10/Subdirect-unions-in-universal-algebra/bams/1183506230.full>.
- [3] John Case and Carl Smith. Comparison of identification criteria for machine inductive inference. *Theoretical Computer Science*, 25(2):193–220, 1983.
- [4] Patrick Cousot and Radhia Cousot. Systematic design of program analysis frameworks. In *Proceedings of the 6th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages*, pages 269–282, 1979. doi: 10.1145/567752.567778. URL <https://cs.nyu.edu/~pcousot/COUSOTpapers/POPL79.shtml>.
- [5] Adnan Darwiche and Pierre Marquis. A knowledge compilation map. *Journal of Artificial Intelligence Research*, 17:229–264, 2002. doi: 10.1613/jair.989.
- [6] E. Mark Gold. Language identification in the limit. *Information and Control*, 10(5):447–474, 1967.
- [7] Eliahu Horwitz, Nitzan Kurer, Jonathan Kahana, Liel Amar, and Yedid Hoshen. We should chart an atlas of all the world’s models. *arXiv preprint arXiv:2503.10633*, 2025.
- [8] Alexander S. Kechris. *Classical Descriptive Set Theory*, volume 156 of *Graduate Texts in Mathematics*. Springer, 1995. doi: 10.1007/978-1-4612-4190-4.
- [9] Miltiadis Kofinas, Boris Knyazev, Yan Zhang, Yunlu Chen, Gertjan J. Burghouts, Efstratios Gavves, Cees G. M. Snoek, and David W. Zhang. Graph neural networks for learning equivariant representations of neural networks. In *International Conference on Learning Representations*, 2024.
- [10] Derek Lim, Haggai Maron, Marc T. Law, Jonathan Lorraine, and James Lucas. Graph metanetworks for processing diverse neural architectures. In *International Conference on Learning Representations*, 2024.
- [11] Francesco Locatello, Dirk Weissenborn, Thomas Unterthiner, Aravindh Mahendran, Georg Heigold, Jakob Uszkoreit, Alexey Dosovitskiy, and Thomas Kipf. Object-centric learning with slot attention. In *Advances in Neural Information Processing Systems*, volume 33, pages 11525–11538, 2020. URL <https://arxiv.org/abs/2006.15055>.
- [12] Tom M. Mitchell. Generalization as search. *Artificial Intelligence*, 18(2):203–226, 1982. doi: 10.1016/0004-3702(82)90040-6.
- [13] Edward F. Moore. Gedanken-experiments on sequential machines. In Claude E. Shannon and John McCarthy, editors, *Automata Studies*, volume 34 of *Annals of Mathematics Studies*, pages 129–153. Princeton University Press, 1956. doi: 10.1515/9781400882618-006.
- [14] National Institute of Standards and Technology. Advanced encryption standard (AES). Technical Report FIPS 197, updated May 9, 2023, U.S. Department of Commerce, 2023.
- [15] Aviv Navon, Aviv Shamsian, Idan Achituve, Ethan Fetaya, Gal Chechik, and Haggai Maron. Equivariant architectures for learning in deep weight spaces. In *International Conference on Machine Learning*, 2023.
- [16] Anil Nerode. Linear automaton transformations. *Proceedings of the American Mathematical Society*, 9(4):541–544, 1958. doi: 10.1090/S0002-9939-1958-0135681-9.
- [17] Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, second edition, 2009. doi: 10.1017/CBO9780511803161.

- 383 [18] Konstantin Schürholt, Diyar Taskiran, Boris Knyazev, Xavier Giró i Nieto, and Damian Borth.
 384 Model zoos: A dataset of diverse populations of neural network models. In *Advances in Neural*
 385 *Information Processing Systems, Datasets and Benchmarks Track*, 2022.
- 386 [19] Thomas Unterthiner, Daniel Keysers, Sylvain Gelly, Olivier Bousquet, and Ilya Tolstikhin.
 387 Predicting neural network accuracy from weights. *arXiv preprint arXiv:2002.11448*, 2020.
- 388 [20] Gail Weiss, Yoav Goldberg, and Eran Yahav. Extracting automata from recurrent neural networks
 389 using queries and counterexamples. In *Proceedings of the 35th International Conference on*
 390 *Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 5247–5256,
 391 2018. URL <https://proceedings.mlr.press/v80/weiss18a.html>.
- 392 [21] Andy Yang, Lena Strobl, David Chiang, and Dana Angluin. Simulating hard attention using soft
 393 attention. *Transactions of the Association for Computational Linguistics*, 14:147–166, 2026.
 394 doi: 10.1162/TACL.a.597. URL <https://aclanthology.org/2026.tacl-1.8/>.
- 395 [22] Allan Zhou, Kaien Yang, Kaylee Burns, Adriano Cardace, Yiding Jiang, Samuel Sokota, J. Zico
 396 Kolter, and Chelsea Finn. Permutation equivariant neural functionals. In *Advances in Neural*
 397 *Information Processing Systems*, volume 36, 2023. URL [https://arxiv.org/abs/2302.](https://arxiv.org/abs/2302.14040)
 398 14040.

A Additional algebraic results

A.1 Continuation equivalence and the minimal-state certificate

We record the details behind the semantic quotient. For every $U, V, H, W \in \mathcal{L}$, if $U \equiv_\rho V$, then

$$\rho((U \cap W) \cap H) = \rho(U \cap (W \cap H)) = \rho(V \cap (W \cap H)) = \rho((V \cap W) \cap H).$$

Thus $U \cap W \equiv_\rho V \cap W$. Applying this twice shows that the meet on equivalence classes is independent of both representatives. Associativity, commutativity, idempotence, and the identity class $[\mathcal{C}]$ descend from intersection. The readout is well-defined because the empty continuation is permitted.

If a correct machine reaches the same state after s, t , determinism gives equal future outputs after every $u \in P^*$. Writing $H = V(u)$ proves $[V(s)] = [V(t)]$. Conversely, the quotient machine computes the target because its state after s is exactly $[V(s)]$. These observations prove both the lower bound and the existence of a smallest implementation. The argument uses the full declared continuation language P^* ; the readout need not identify individual candidates.

A finite certificate for a proposed implementation consists of its reachable transition table, a map δ into $\mathcal{R}$, and the three equations in (3). Verifying the equations for every reachable state and every input symbol proves correctness for all histories by induction. With N reachable states and m input symbols, the transition part has Nm equalities. This is an all-horizon certificate rather than a test on histories up to a chosen length.

A.2 Idempotent lifting in finite commutative monoids

For completeness, let x belong to a finite monoid. There exist positive integers a, b with $x^a = x^{a+b}$. Hence $x^t = x^{t+b}$ for all $t \geq a$. Choose a multiple r of b with $r \geq a$. The exponents $r, 2r$ differ by a multiple of b , so $x^{2r} = x^r$. Thus x^r is idempotent.

In a commutative monoid, the product of two idempotents is idempotent. On its idempotents, $e \leq f$ if $ef = e$ defines a partial order. Reflexivity is idempotence; antisymmetry follows from commutativity; and if $ef = e$, $fg = f$, then $eg = (ef)g = e(fg) = ef = e$. The product ef is the greatest lower bound of e, f , and the monoid identity is top.

Let $\pi : S \rightarrow R$ be a surjective monoid homomorphism from a finite commutative monoid to a meet-semilattice with top. Since $\pi(x^r) = \pi(x)$ for every $r \geq 1$, the idempotents of S still map onto R . Given $r_1 < \dots < r_h$, choose idempotent preimages e_i and let $z_i = e_i \dots e_h$. Their images are r_i , and $z_i z_{i+1} = z_i$. This proves the chain-lifting assertion used in theorem 3.

The chain height of a product of finite posets satisfies

$$\text{ht} \left(\prod_{j=1}^d Q_j \right) \leq 1 + \sum_{j=1}^d (\text{ht}(Q_j) - 1).$$

Along a strict chain, coordinate j can change strictly at most $\text{ht}(Q_j) - 1$ times, and every strict step changes some coordinate. For finite posets the upper bound is attained by interleaving longest chains in the factors. Restricting to a subposet cannot increase height. These facts complete the proof for pooled implementations with arbitrary redundant states.

A.3 Subdirect representations and residual size

Proposition 12. *Let $\alpha : R \rightarrow A$ and $\beta : R \rightarrow G$ be meet homomorphisms on a finite meet-semilattice R , with kernels α, β . The paired representation determines every state of R if and only if $\alpha \cap \beta = \Delta_R$. If R is the minimal inference state, a joint readout correct after every continuation exists if and only if this condition holds.*

Proof. Two states have the same pair precisely when they are equivalent in both kernels. Hence the kernel condition is equivalent to injectivity. Injection defines the decoded state and therefore a correct readout. Conversely, if distinct minimal states share a pair, they continue to share it after every meet update because both views are homomorphisms. Minimality supplies a distinguishing continuation, contradicting correctness. $\square$

Every homomorphic image of R is isomorphic to its kernel quotient. Conversely, each congruence quotient supplies a compositional view. The set in (13) is nonempty because the equality congruence is feasible. This proves that the minimum is exact, not merely a lower bound. For $\alpha' \subseteq \alpha$, every β feasible for α is feasible for α' , proving monotonicity under structural refinement.

For a static residual $g : R \rightarrow G$ with no compositional requirement, injectivity of (a, g) requires at least $|F|$ residual values on each structural fibre F . Let m be the largest fibre size and inject every fibre separately into $[m]$. The resulting residual attains m , proving (14). It encodes a one-shot reconstruction of the state.

For a chain, every equivalence relation with interval blocks is a meet congruence: if $x \sim y$, taking minimum with z either preserves the pair, maps both to z , or produces two points between x, y . The interval property gives equivalence in each case. Therefore every choice of adjacent cuts defines a congruence and all chain congruences arise this way. If the structural view cuts a set A of boundaries, the residual cuts must contain its complement. They require $n - |A| = n - a + 1$ blocks, and no additional cuts are necessary.

The Boolean coordinate law follows without requiring a normalized residual. Fix a structural fibre and choose a reaching history for each of its 2^{k-s} states. Correctness and proposition 2 require distinct joint implementation states. Since their structural values agree, the residual values are pairwise distinct. This lower bound applies to any finite independently pooled residual. The complementary-coordinate projection attains it.

A.4 A presentation with many survivors but one inference state

Let $\mathcal{C} = \{c_0, c_1, \dots, c_k\}$, let observations independently remove any of $c_1, \dots, c_k$, and let no observation remove c_0 . The reachable version spaces are all subsets containing c_0 , so $|\mathcal{L}| = 2^k$. First-fit selection in the displayed order always returns c_0 , including after every continuation. Hence $|\mathcal{R}| = 1$. This example makes the readout dependence of the lower bound explicit. Counting reachable survivor sets is sharp precisely when the continuation equivalence separates them.

B Extraction and finite abstractions

B.1 Operational reconstruction of the meet

Proposition 13 (Recovering the inference algebra). *Let $(Q, q_0, (t_p)_{p \in P}, o)$ be a finite reachable deterministic Moore machine. Suppose the transition maps commute and are idempotent:*

$$t_p t_q = t_q t_p, \quad t_p^2 = t_p \quad (p, q \in P).$$

The operation $r \wedge s = t_{w_s}(r)$, for an access word w_s of s , is a meet with top q_0 satisfying

$$t_p(q) = q \wedge t_p(q_0),$$

and any meet operation with top q_0 satisfying this identity equals it. Uniqueness uses reachability and the identity alone; commutativity and idempotence enter through existence. The operation can be reconstructed from the transition table. Its version-space realization on candidate set Q has consistency sets $K(p) = \downarrow t_p(q_0)$ and survivor state $\downarrow q_s$ after history s .

Proof. For a word w , write t_w for its transition map. Commutativity and idempotence of the generators imply these properties for all t_w . For each state r , choose an access word w_r with $t_{w_r}(q_0) = r$. Define $r \wedge s = t_{w_s}(r)$. If w'_s is another access word for s , then, using an access word for r ,

$$t_{w_s}(r) = t_{w_s} t_{w_r}(q_0) = t_{w_r}(s) = t_{w_r} t_{w'_s}(q_0) = t_{w'_s}(r).$$

Thus the definition is independent of the access word. It is commutative by the same calculation. The concatenation $w_s w_t$ is an access word for $s \wedge t$, which proves associativity. Idempotence follows from $t_{w_r}^2 = t_{w_r}$, and the empty word gives identity q_0 . Therefore the operation is a meet with top.

Taking the one-letter access word for $t_p(q_0)$ gives the transition identity. In any other meet representation with that identity, induction on an access word forces $r \wedge s = t_{w_s}(r)$, using only reachability and the identity; this proves uniqueness. Finally,

$$\downarrow r \cap \downarrow s = \downarrow(r \wedge s), \quad \downarrow q_0 = Q.$$

487 Induction on observations gives the stated survivor sets. The map $r \mapsto \downarrow r$ is injective, and the readout
 488 is $\rho(\downarrow r) = o(r)$. $\square$

489 For a minimal machine whose output on histories is invariant to permutation and repetition, the
 490 transition maps necessarily commute and are idempotent. Indeed, applying pq or qp after any reaching
 491 history, followed by any continuation, gives equal outputs; minimality makes the resulting states
 492 equal. The same argument compares pp with p . Conversely, commuting idempotent transitions
 493 imply both invariances directly. Consequently the condition in proposition 13 can be tested from the
 494 extracted minimal machine, without presupposing a model population.

495 The resulting model is operational. It identifies the distinctions implemented by the artifact under
 496 its observation interface. Connecting those states to external hypotheses, causal mechanisms, or
 497 object identities requires a semantic interpretation of the observations and readout. When such an
 498 interpretation is supplied by (1), proposition 2 makes the correspondence unique on reachable states.

499 B.2 Partition refinement and certificates

500 Let N be the number of reachable states. Start with equivalence of output labels and refine by

$$q \sim_{r+1} q' \iff o(q) = o(q') \text{ and } t_p(q) \sim_r t_p(q') \text{ for all } p \in P.$$

501 Induction shows that $\sim_r$ is agreement on all continuations of length at most r . Each strict refinement
 502 increases the number of blocks, so stabilization occurs after at most $N - 1$ strict refinements. At
 503 stabilization the relation is a congruence for every transition. Induction on arbitrary continuations
 504 then proves agreement at every length. For two different blocks, the refinement history yields a
 505 finite distinguishing word. Thus a complete extraction certificate consists of the transition table, the
 506 quotient projection, and distinguishing continuations between distinct blocks.

507 For a finite input alphabet, testing commutativity and idempotence requires finitely many state equali-
 508 ties. Once proposition 13 applies, its access-word construction gives the meet table. Enumerating
 509 finite set partitions and testing meet compatibility enumerates every meet congruence. Evaluating
 510 (13) is therefore an exact finite procedure on an extracted artifact. These quantities are invariant under
 511 isomorphisms of the extracted machine.

512 B.3 Continuous-state artifacts with finite invariant regions

513 Exact extraction is not confined to one-hot activations. Let $Z \subseteq \mathbb{R}^d$ be an invariant state set and let
 514 $B_1, \dots, B_N$ be a finite partition of Z . Assume that each input p maps every region into one region
 515 and that the readout is constant on each region:

$$F_\theta(B_i, p) \subseteq B_{T(i,p)}, \quad o_\theta(z) = o_i \quad (z \in B_i), \quad z_0 \in B_{i_0}. \quad (30)$$

516 **Proposition 14** (Finite-region extraction). *The finite machine (T, i_0, o) has exactly the same output*
 517 *on every history as the continuous-state artifact. If Z and its state regions have rational semialgebraic*
 518 *descriptions, the initial state and input encodings are rational, output labels have rational encodings,*
 519 *and the transition and readout consist of rational affine, ReLU, and deterministically tie-broken*
 520 *hard-attention operations, all conditions in (30), including coverage and disjointness, are decidable.*

521 *Proof.* The initial state belongs to the declared initial region. If the current state belongs to B_i ,
 522 the inclusion in (30) puts its successor in $B_{T(i,p)}$; the readout equality then proves the simulation.
 523 Induction gives equality for every history.

524 Affine and ReLU operations have semialgebraic graphs. A selected argmax index j is described by
 525 the finite conjunction $s_j > s_i$ for $i < j$ and $s_j \geq s_i$ for $i > j$. Hence a fixed hard-attention network
 526 has a semialgebraic graph. Each inclusion or readout equality in (30) is a first-order sentence over the
 527 ordered real field with rational coefficients. Real quantifier elimination decides it [1]. Coverage and
 528 disjointness are sentences of the same form. $\square$

529 The proposition gives a certificate checker for a supplied abstraction. The same partition-refinement
 530 procedure extracts its minimal machine. Region labels need not be the final semantic states: mini-
 531 mization removes redundant distinctions between regions.

C A literal attention–ReLU compiler

C.1 Leaf maps and transition maps

For a map $f : [m] \rightarrow [q]$, set the i th column of H_f to $e_{f(i)}$. Then $H_f e_i = e_{f(i)}$. A ReLU layer with identity input weights is the identity on nonnegative one-hot inputs, so a standard two-affine-layer feedforward block implements the table exactly. For a map on two finite inputs, the units $\text{ReLU}(z_i + p_j - 1)$ in (19) are zero except at the unique active pair. Their output matrix therefore gives the exact transition table. The construction uses Nm hidden units and rational weights in $\{-1, 0, 1\}$ for its conjunction and one-hot output stages.

A leaf table is an operational parameter matrix. The extractor obtains it by evaluating the leaf on every basis input, whether or not the implementation’s hidden units have been permuted or reorganized. The same procedure extracts a transition table from any implementation satisfying the finite interface, not just from the displayed compiler.

C.2 Tuple formation and finite dataflow

To combine r already computed symbols, use one fixed-mask attention head per input to retrieve its one-hot vector $v^{(1)}, \dots, v^{(r)}$ into disjoint scratch blocks. For a tuple $(i_1, \dots, i_r)$, the hidden unit

$$h_{i_1, \dots, i_r} = \text{ReLU}\left(\sum_{a=1}^r v_{i_a}^{(a)} - (r - 1)\right) \quad (31)$$

is one exactly at the active tuple and zero elsewhere. An output matrix therefore implements any finite map of the retrieved tuple. Shared node-specific dispatch uses the unit $\text{ReLU}(u_v + \sum_a v_{i_a}^{(a)} - r)$ instead. A topological schedule of these blocks implements every finite acyclic dataflow program. Fixed-mask reads have a single allowed key, so their outputs are identical for hard and soft attention. In particular, the tuple construction composes the forest outputs and implements parent retrieval and structural equations in appendix F; no external tuple-encoding oracle is needed.

C.3 Hard attention gates

Write each routing score as

$$s_j(x) = \alpha_j^\top \phi(x) + \beta_j = \langle (\phi(x), 1), (\alpha_j, \beta_j) \rangle.$$

The child tokens carry key (α_j, β_j) and value $e_{\llbracket c_j \rrbracket(x)}$. The parent query is $(\phi(x), 1)$. The usual attention scale can be absorbed into the keys. With the declared deterministic tie rule, hard attention returns the value of $j^*(x)$. Multiplication by H_g gives the node’s value. Induction from leaves proves the semantic identity at every node.

The extractor reads the effective query–key score coefficients, the wiring or mask, and the finite table induced by the feedforward block. Different factorizations into query and key matrices can give the same scores; the extracted code is determined semantically rather than by selecting one parameter factorization. Child reorderings preserve the extracted semantics when the tie convention is transported with the ordering. The equality in (18) concerns the decoded code, not syntactic equality with its source.

C.4 Shared positionwise feedforward maps and residual paths

We give a standard masked-transformer realization of the forest dataflow. Unfold a finite directed acyclic graph into a finite forest when a node has multiple parents. Use one token for each resulting node. Embed the finite node-output alphabets into a common tagged alphabet, so every payload uses a common coordinate block. The token coordinates contain: a fixed one-hot node identifier; protected input features $(\phi(x), 1)$; the node’s key for its parent; a one-hot payload; and a scratch block initialized to zero. The identifiers and keys are operational inputs to shared dispatch and routing.

Process nodes by height above the leaves. In a layer, the mask of each active parent exposes exactly its children. Inactive tokens attend to themselves. Shared query and key projections select the protected query and key blocks. The shared value projection reads the payload; the output projection writes

into scratch. Because scratch was zero, the attention residual leaves there exactly the attended value. Protected coordinates and the old payload remain unchanged.

For hard attention, scratch is one-hot. Let u_v be the identifier of node v , and let w_j be the scratch symbol coordinate. The shared hidden unit

$$h_{vj} = \text{ReLU}(u_v + w_j - 1)$$

is active precisely for token v and its selected symbol j . An output matrix maps it to the payload required by that node’s finite map. Inactive nodes instead retain their old payload using the same conjunction construction with old-payload coordinates. The feedforward residual update subtracts the old payload and adds the desired one, and subtracts scratch to reset it to zero. Subtraction of a nonnegative coordinate is implemented by a negative output weight on its ReLU. All protected coordinates receive zero update. Thus one shared positionwise feedforward map implements every node-specific operation at that layer.

A preliminary finite-map layer initializes leaf payloads from the one-hot input and initializes internal payloads to a fixed dummy symbol. Hence every payload is a valid one-hot vector from the start. Induction over heights proves that active children have already acquired the correct symbol and that completed nodes retain it. A forest of depth D needs at most D routing layers after this initialization, followed by the declared root readout. This construction fixes all masks, projections, residual paths, and feedforward dispatch; it introduces no unused field encoding the original program identifier.

C.5 Finite-symbol correction within the shared block

Correction can also be implemented in a shared single-hidden-layer feedforward block. For a token identifier $u_v \in \{0, 1\}$ and scratch coordinate $w_j \in [0, 1]$, define

$$h_{vj}^\eta = \frac{\text{ReLU}(w_j + u_v - 1 - \eta) - \text{ReLU}(w_j + u_v - 2 + \eta)}{1 - 2\eta}. \quad (32)$$

If $u_v = 1$, this equals $R_\eta(w_j)$. If $u_v = 0$, both ReLUs vanish because $w_j \leq 1$. The output matrix can therefore apply the node-specific finite map to the corrected symbol. The payload and scratch overwrite construction is unchanged. Whenever (22) holds, the outgoing payload is exactly one-hot at each layer. This gives a literal implementation of theorem 9, with shared parameters across token positions.

For real inputs beyond a finite input alphabet, the routing induction remains valid on any domain on which leaf maps have the stated implementation and the gate conditions hold. In the finite construction, all feature values and leaf maps are part of the specified input embedding, so every operation used in extraction is explicitly available.

D Hardening, perturbations, and observability

D.1 Sharpness of the symbol-mass bound

Let the total unnormalized weight of the target symbol be G and that of all other symbols be B . Then $\varepsilon_{\neq y^*} = B/(G + B)$. Under the hypotheses of theorem 9, $G \geq m_* e^{s_*/T}$ and $B \leq k_- e^{(s_* - \delta)/T}$, yielding (23). Equality is attained when the m_* good branches all have score s_* , all k_- bad branches have score $s_* - \delta$, and no other branches are present. Solving $r/(1 + r) \leq \eta$ gives $r \leq \eta/(1 - \eta)$ and the unique-winner temperature bound (24).

The local criterion in (22) is necessary and sufficient for correction to the selected one-hot symbol. A later finite map can collapse several symbols and thereby permit further equivalences; the criterion is applied before that map. It already accounts for branch equivalence at the symbol level. In particular, if every branch returns the same symbol, $\varepsilon_{\neq y^*} = 0$ for all temperatures, regardless of routing ties.

D.2 A perturbation certificate

Corollary 15 (Numerical robustness of correction). *Let the ideal attention mixture have wrong-symbol mass at most $\eta - \xi$ for some $\xi > 0$. If its numerical realization differs coordinatewise by at most ξ , then applying R_η still returns the exact selected one-hot symbol.*

619 *Proof.* The selected coordinate is at least $1 - \eta + \xi$ before perturbation and at least $1 - \eta$ afterward.
 620 Every other coordinate is at most $\eta - \xi$ before perturbation and at most η afterward. The correction
 621 function is constant on these two regions. $\square$

622 If each score is perturbed by at most ζ and a unique winning score has gap $\delta > 2\zeta$, the winner is
 623 unchanged and the perturbed gap is at least $\delta - 2\zeta$. Substituting this gap into (23), followed by
 624 corollary 15, gives a score-and-value perturbation certificate. The same-symbol version permits route
 625 changes within a group whose values coincide; it suffices to preserve a positive gap from that group's
 626 largest score to every different-symbol branch.

627 D.3 Approximate hardening without symbolic correction

628 For bounded expert values $\|v_j\| \leq B$ and a unique winning score with gap δ , put $r = (k - 1)e^{-\delta/T}$.
 629 Directly,

$$\left\| \sum_j p_j(T) v_j - v_{j^*} \right\| \leq 2B \frac{r}{1+r} \leq 2B(k-1)e^{-\delta/T}. \quad (33)$$

630 For compositions $F_L^T \circ \dots \circ F_1^T$ and $F_L^0 \circ \dots \circ F_1^0$, assume that a common region contains all compared
 631 trajectories, each F_ℓ^T is L_ℓ -Lipschitz there, and $\sup_z \|F_\ell^T(z) - F_\ell^0(z)\| \leq \varepsilon_\ell(T)$. Inserting and
 632 subtracting $F_\ell^T(z_{\ell-1}^0)$ yields

$$e_\ell \leq L_\ell e_{\ell-1} + \varepsilon_\ell(T), \quad e_L \leq \sum_{\ell=1}^L \varepsilon_\ell(T) \prod_{j=\ell+1}^L L_j.$$

633 Uniform score separation is imposed on the region where the local hard/soft bound is used. Al-
 634 ternatively, a stepwise perturbation argument certifies preservation of each gate's winning symbol.
 635 These hypotheses expose the dependence of approximate hardening on downstream stability; exact
 636 symbolic correction removes that accumulation by restoring a codebook element after each gate.

637 Let $\widehat{L}_T(\theta)$ and $\widehat{L}_0(\theta)$ be empirical losses on a fixed sample. If $|\widehat{L}_T(\theta) - \widehat{L}_0(\theta)| \leq \varepsilon$ for every θ in a
 638 declared parameter set, then their approximate-fit sets obey

$$\{\theta : \widehat{L}_T(\theta) \leq a - \varepsilon\} \subseteq \{\theta : \widehat{L}_0(\theta) \leq a\} \subseteq \{\theta : \widehat{L}_T(\theta) \leq a + \varepsilon\}.$$

639 This follows by adding the uniform loss discrepancy. Under exact symbolic preservation and a loss
 640 on the resulting symbols, the two sets are equal for every threshold.

641 D.4 Artifact observations and their sufficiency

642 For maps $O : \Theta \rightarrow \mathcal{O}$, $A : \Theta \rightarrow \mathcal{A}$, and $Q : \Theta \rightarrow \mathcal{Q}$, an exact inference function on the image of
 643 (O, A) exists if and only if Q is constant on its fibres. Necessity follows by applying the inference
 644 function to equal observation pairs. For sufficiency, assign to each pair in the image the common
 645 value of Q on its fibre. This defines a function with the required factorization and proves (26).

646 If $O_0 = h \circ O_T$, equality of O_T implies equality of O_0 , proving (25). This is an observational
 647 comparison, with a declared coarsening map. It does not equate the activations of independently
 648 run soft and hard networks when their upstream states change. Theorem 9 supplies the separate
 649 execution-equivalence theorem for the compiled symbolic interfaces.

650 E Sound exchange between representations

651 Let $\mathcal{A}, \mathcal{G} \subseteq \mathcal{P}(\mathcal{C})$ be families containing $\mathcal{C}$ and closed under arbitrary intersections. Since $\mathcal{C}$ is finite,
 652 it suffices to require closure under finite intersections. Define

$$\text{cl}_{\mathcal{A}}(V) = \bigcap \{A \in \mathcal{A} : V \subseteq A\}, \quad \text{cl}_{\mathcal{G}}(V) = \bigcap \{G \in \mathcal{G} : V \subseteq G\}.$$

653 A structural state A and residual state G jointly denote $A \cap G$. Their reduced product [4] is

$$\text{red}(A, G) = (\text{cl}_{\mathcal{A}}(A \cap G), \text{cl}_{\mathcal{G}}(A \cap G)).$$

654 **Proposition 16.** *Reduction is componentwise contracting, preserves the joint version space exactly,*
655 *and is idempotent.*

656 *Proof.* Write $V = A \cap G$, $A' = \text{cl}_A(V)$, and $G' = \text{cl}_G(V)$. Since A and G are available overapprox-
657 imations, $V \subseteq A' \subseteq A$ and $V \subseteq G' \subseteq G$. Hence $V \subseteq A' \cap G' \subseteq A \cap G = V$. Repeating reduction
658 uses the same V and therefore produces the same pair. $\square$

659 The representation in one component can sharpen the other by making consequences of their joint
660 information expressible there. No additional external observation is needed for this transfer. The joint
661 set of possible models is unchanged, so the operation preserves evidence.

662 F Model queries and causal interpretation

663 F.1 Query certificates and adaptive experiments

664 Fix $c^* \in V$. After truthful outcomes on E , the survivor set is

$$V_E = \{c \in V : D(c, e) = D(c^*, e) \text{ for all } e \in E\}.$$

665 The true candidate belongs to this set. Intersecting with B_q gives $B_q \setminus \bigcup_{e \in E} A_e$, proving propo-
666 sition 10. For a stored experiment set E_0 , an additional set E_1 certifies the query exactly when
667 its kill sets cover $B_q \setminus \bigcup_{e \in E_0} A_e$. Model identification is the special case $q(c) = c$, or a semantic
668 equivalence class of c when extensional duplicates are identified.

669 The result is target-relative. A policy choosing experiments before learning which model is true is
670 represented by a decision tree whose branches are outcomes. Every root-to-leaf path must satisfy the
671 same covering condition for each true model reaching that leaf. Thus the proposition characterizes
672 the certificates required at policy leaves; optimizing the entire adaptive policy additionally chooses
673 how those certificates share prefixes.

674 F.2 Compiling a finite acyclic structural causal model

675 Let the endogenous variables $X_1, \dots, X_d$ and exogenous variable U have finite alphabets. For each
676 model c , let its graph be acyclic, with structural maps

$$X_i = f_{c,i}(X_{\text{pa}_c(i)}, U).$$

677 For a fixed intervention $I = \text{do}(X_J = a_J)$ and exogenous input u , evaluate the variables in a
678 topological order, using the constant a_i for $i \in J$ and the structural table otherwise. Each operation is
679 a finite map. Retrieve its already evaluated parent symbols and the exogenous symbol with fixed-mask
680 attention heads. The conjunction units in (31) form the tuple basis vector, and the table matrix applies
681 its structural map. An intervention-controlled finite map selects the constant branch or the structural
682 branch. The forest compiler realizes these operations by attention and ReLU maps.

683 **Proposition 17** (Intervention preservation). *For every model c , intervention I , and exogenous input*
684 *u , the compiled program returns the unique solution of the intervened structural equations. For a*
685 *specified finite exogenous distribution, weighted aggregation gives the exact intervention distribution.*

686 *Proof.* At the first variable in a topological order, the program applies either its intervention constant
687 or the structural map with exogenous input u . This agrees with the intervened equation. Suppose
688 the program is correct for all preceding variables. Every parent of the next variable precedes it, so
689 the retrieved tuple is correct; applying its selected table proves the next value correct. Induction
690 gives the unique solution. Summing its one-hot output distribution against the specified exogenous
691 probabilities gives the exact interventional law. $\square$

692 To apply the deterministic certificate theorem, an experiment can report a controlled-exogenous
693 outcome or the exact observable distribution under an intervention. The output of such an experiment
694 is a deterministic function of the candidate model. Observations from a stochastic experiment require
695 their own statistical consistency map; the survivor and continuation constructions then use that
696 declared map.

697 For an explicit separation, let U be uniform on $\{0, 1\}$ and consider

$$c_1 : X = U, Y = X, \quad c_2 : X = U, Y = U.$$

698 Both give the same observational distribution: probability $1/2$ each on $(0, 0)$ and $(1, 1)$. Under
 699 $\text{do}(X = 0)$, the first model has $Y = 0$ surely, while the second retains uniform Y . The query
 700 $\Pr(Y = 1 \mid \text{do}(X = 0))$ is therefore 0 or $1/2$. Observational distribution access does not separate
 701 the pair, while that intervention distribution does. Their structural tables, or a compiled program with
 702 the same intervention semantics, also separate the query by (26).

703 G Population memory bounds and routed model classes

704 **Proposition 18** (Table reconstruction). *Suppose a fixed deterministic decoder reconstructs every*
 705 *function $f : [N] \rightarrow [q]$ from a target-dependent B -bit artifact and at most t adaptive table queries.*
 706 *Then*

$$B + t \log_2 q \geq N \log_2 q.$$

707 *For a population of all permutations of $[N]$, with $0 \leq t \leq N$, the corresponding bound is*

$$B \geq \log_2((N - t)!).$$

708 *Proof.* Fix the artifact. Its adaptive query tree has at most q^t leaves after padding shorter transcripts
 709 to depth t . The artifact and answer transcript determine all queries and the reconstructed table. Two
 710 different tables cannot have the same pair, since the decoder is exact. There are at most $2^B q^t$ pairs
 711 and q^N tables, giving the first inequality.

712 For permutations, repeated queries supply no new information and can be omitted. Pad to t distinct
 713 queries. The possible answer counts at successive levels are at most $N, N - 1, \dots, N - t + 1$,
 714 independent of how query locations are chosen. There are at most $N!/(N - t)!$ leaves for each
 715 artifact. Comparing the $2^B N!/(N - t)!$ pairs with the $N!$ targets proves the second inequality. $\square$

716 All target-dependent information, including a varying architecture, routing parameters, learned
 717 encoders, and finite tables, belongs in the B -bit description. The bound concerns a population with a
 718 fixed decoder, not the shortest program of a particular function. A restricted algebraic population
 719 can have a shorter description. This is why a structured S-box implementation and an arbitrary
 720 permutation table belong to different counting problems.

721 G.1 Selection and execution have different statistical complexity

722 Let $\mathcal{F}$ consist of routed functions $f(x) = h_{r(x)}(x)$, with routing class $\mathcal{H}_{\text{route}}$ and expert classes $\mathcal{H}_i$,
 723 all with finite output alphabets. For a finite sample S , let $S_i(r)$ be the sample positions routed to
 724 expert i . Then

$$|\mathcal{F} \upharpoonright_S| \leq \sum_{r \in \mathcal{H}_{\text{route}} \upharpoonright_S} \prod_i |\mathcal{H}_i \upharpoonright_{S_i(r)}|. \quad (34)$$

725 For each realized routing pattern, the output restriction is determined by the expert restrictions on
 726 their assigned positions. Multiplying counts bounds that pattern; summing over patterns gives (34).
 727 Fixed experts contribute factors of one. When expert contents also vary, their restrictions contribute
 728 explicitly. Thus an economical selection representation does not erase the variability of the models it
 729 executes.

730 H Hardening and measurability

731 Discontinuity of argmax is not, by itself, a failure of measurable inference. A deterministic finite
 732 argmax of measurable scores is measurable: each selected-index event is a finite conjunction of
 733 measurable comparisons, with strict comparisons for earlier indices to implement the tie rule. Finite
 734 compositions with measurable feedforward maps remain measurable.

735 **Proposition 19** (Finite parameter families). *For a finite parameter set Θ and measurable losses $\ell_\theta(z)$*
736 *with defined finite expectations $L(\theta)$, the uniform-deviation event*

$$\left\{ (z_1, \dots, z_m) : \max_{\theta \in \Theta} \left| L(\theta) - \frac{1}{m} \sum_{i=1}^m \ell_\theta(z_i) \right| > \varepsilon \right\}$$

737 *is measurable. This holds for both soft and deterministically tie-broken hard attention whenever their*
738 *individual losses are measurable.*

739 *Proof.* Each empirical loss is measurable in the sample. The maximum is finite, so its threshold event
740 is a finite union of measurable events. $\square$

741 A second statement applies to real-valued parameters. Fix a finite architecture of affine, ReLU, and
742 hard-attention operations, a semialgebraic parameter set $\Theta \subseteq \mathbb{R}^p$, and a semialgebraic loss. If the
743 reference distribution has finite support, then the population loss is semialgebraic in the parameters.
744 The graph of the network is semialgebraic by the same argmax description used in proposition 14. The
745 joint strict-deviation set in parameter and sample variables is therefore semialgebraic. Its projection
746 to sample variables is semialgebraic by the Tarski–Seidenberg theorem, hence Borel [1].

747 For general parameter spaces and joint families, projection raises a separate descriptive-set-theoretic
748 issue: projections of Borel sets need not be Borel [8]. The relevant object is the parameter-uniform
749 event, not an individual network evaluation. The preservation and observation-coarsening results in
750 section 5 identify how hardening changes computation and identifiability without identifying either
751 change with a loss of measurability.

752 I Reproducibility and executed checks

753 All results in this paper are verified in Lean 4 with Mathlib. The accompanying implementation
754 consists of four Python modules and a machine-readable result file. Running

755 `python checks/verify.py`

756 regenerates every reported count. The recorded environment is Python 3.13.5 with NumPy 2.3.5.
757 Integer checks use exact integer arithmetic. Float64 is used for dot products and softmax evaluation
758 in the finite attention constructions. No fitted parameters, external datasets, or randomized train/test
759 splits enter the enumerations.

760 **Inference quotients.** For candidate sizes $k = 1, 2, 3$, enumerate every family of bit masks contain-
761 ing the full mask and closed under bitwise intersection. There are 70 such labeled families in total.
762 For each family, enumerate every Boolean readout and add the first-surviving-index readout, with a
763 separate inconsistency output. This gives 2,120 cases. Direct equivalence by all meets is compared
764 against iterative transition refinement; readout descent and meet compatibility are also checked.

765 **Congruences and exact residual minima.** Enumerate every set partition of each chain of size
766 one through seven, retaining exactly those compatible with minimum. There are 2^{n-1} at size n and
767 127 structural quotients across the range. For each, search all residual congruences for the smallest
768 jointly injective representation. The result is $n - a + 1$ in every case. Repeat for Boolean stores
769 on one through three coordinates and every initial coordinate projection; the minima are 2^{k-s} . The
770 four-chain parity residual is explicitly checked to be injective jointly with the structural view but
771 incompatible with minimum.

772 **Non-idempotent pooling.** Enumerate every commutative multiplication table with identity labeled
773 zero on carriers of sizes one through four, and retain the associative tables. The counts are 1, 2, 9, 94,
774 totaling 106 labeled monoids; 93 are non-idempotent. For every element, a factorial power is checked
775 to be idempotent. Enumerate every surjective homomorphism from each monoid to a chain of size no
776 larger than its carrier. All 329 maps are checked to remain surjective on idempotents and satisfy the
777 height bound. These are labeled tables, not isomorphism classes.

778 **Neural machine extraction.** Compile each of the 2,120 inference machines by (19). Apply a
779 hidden-unit permutation using seed 20260907, transporting both adjacent weight matrices. Extract
780 the transition and readout tables from the resulting network and minimize them. All 72,911 state-
781 input evaluations reproduce the specified transition, and every extracted quotient equals the direct
782 continuation equivalence.

783 **Attention forests.** The input set is $\{0, \dots, 7\}$. The four base maps are constant zero, constant one,
784 parity, and complemented parity. Enumerate all one-gate combinations of these maps over thresholds
785 1.5, 3.5, 5.5, with either the identity or complemented output map. This gives 96 one-gate codes, in
786 addition to four leaves. Finally enumerate all $4^4 = 256$ base-map assignments to the four leaves
787 of a balanced depth-two tree whose thresholds are 1.5, 3.5, 5.5. The total is 356 network artifacts,
788 realizing all 256 Boolean functions on eight inputs.

789 Routing scores at threshold t are $t - x$ and $x - t$. Their minimum gap on the integer input set is
790 one. The feedforward correction uses $\eta = 1/4$ and then applies the chosen finite map. Extracted
791 routing coefficients, leaf tables, and output maps reproduce the code on all 2,848 network-input
792 pairs. Softmax execution at temperatures 0.1, 0.25, and 0.5 gives 8,544 further pairs. The maximum
793 recorded absolute error from the target one-hot output is zero in float64. These equalities check the
794 compiled finite-symbol construction; the mathematical guarantee is theorem 9.

795 **Literal shared-block execution.** Compile the same 356 forests into token matrices with protected
796 node identifiers, input features, parent keys, payloads, and scratch blocks. Execute the actual shared
797 query/key/value/output projections, masked attention, and shared ReLU feedforward matrices with
798 both residual paths. For each artifact and each of the eight inputs, run hard attention and temperatures
799 0.1, 0.25, and 0.5. All 11,392 runs return the target symbol with zero maximum recorded absolute
800 error. After every layer, protected coordinates are unchanged, payloads are one-hot, and scratch is
801 zero; the largest recorded error in these invariants is also zero. The forward evaluator receives only
802 matrices, masks, and the input, not the source program.

803 **Same-symbol ties and coupled minimum.** Enumerate every three-branch score vector in
804 $\{-1, 0, 1\}^3$, every binary assignment of branch values, and temperatures 0.1, 0.25, 0.5, 1, 2. The
805 1,080 cases verify the necessary-and-sufficient wrong-symbol-mass criterion; 210 involve tied
806 maximum-score branches with identical symbols. Separately, allocate chain cuts for sizes two
807 through 32 and alphabet sizes two through eight, giving 217 sharp independent-width constructions.
808 The coupled digit-code attention minimum is verified for chain sizes two through 64 and alphabets
809 2, 3, 4, 8, covering 357,756 input pairs.